

VYSOKÉ UČENÍ TECHNICKÉ V BRNĚ

BRNO UNIVERSITY OF TECHNOLOGY

FAKULTA INFORMAČNÍCH TECHNOLOGIÍ
ÚSTAV INFORMAČNÍCH SYSTÉMŮ

FACULTY OF INFORMATION TECHNOLOGY
DEPARTMENT OF INFORMATION SYSTEMS

MODELLING NEW NETWORK ARCHITECTURES IN OMNET++

BAKALÁŘSKÁ PRÁCE
BACHELOR'S THESIS

AUTOR PRÁCE
AUTHOR

TOMÁŠ HYKEL

BRNO 2015



VYSOKÉ UČENÍ TECHNICKÉ V BRNĚ
BRNO UNIVERSITY OF TECHNOLOGY



FAKULTA INFORMAČNÍCH TECHNOLOGIÍ
ÚSTAV INFORMAČNÍCH SYSTÉMŮ

FACULTY OF INFORMATION TECHNOLOGY
DEPARTMENT OF INFORMATION SYSTEMS

MODELOVÁNÍ NOVÝCH SÍŤOVÝCH ARCHITEKTUR V OMNET++

MODELLING NEW NETWORK ARCHITECTURES IN OMNET++

BAKALÁŘSKÁ PRÁCE

BACHELOR'S THESIS

AUTOR PRÁCE

AUTHOR

TOMÁŠ HYKEL

VEDOUCÍ PRÁCE

SUPERVISOR

Ing. MARCEL MAREK

BRNO 2015

Abstrakt

Abstract

Klíčová slova

Keywords

Citace

Tomáš Hykel: Modelling New Network Architectures in OMNeT++, bakalářská práce, Brno, FIT VUT v Brně, 2015

Modelling New Network Architectures in OMNeT++

Prohlášení

Prohlašuji, že jsem tuto bakalářskou práci vypracoval sám pod vedením Ing. Marcela Marka

.....

Tomáš Hykel
March 12, 2015

Poděkování

© Tomáš Hykel, 2015.

Tato práce vznikla jako školní dílo na Vysokém učení technickém v Brně, Fakultě informačních technologií. Práce je chráněna autorským zákonem a její užití bez udělení oprávnění autorem je nezákonné, s výjimkou zákonem definovaných případů.

Contents

1	Introduction	3
1.1	Goals	3
1.2	Thesis Structure	3
2	Problems of the current Internet	4
2.1	Node naming	4
2.2	Lack of Multihoming	5
2.2.1	SCTP	5
2.2.2	BGP Multihoming	5
2.2.3	Multipath TCP	5
2.3	Lack of Mobility	5
2.4	Lack of Multicast	6
2.5	Weak security	6
2.6	Router Table Size Growth	6
3	Overview of Alternative Architectures	8
3.1	Design Approaches	8
3.2	Research Initiatives	8
3.2.1	NewArch Project	8
3.2.2	FIRE	8
3.2.3	FIBRE	8
3.2.4	Future Internet Architecture	8
3.3	Named Data Networking	9
3.3.1	Premise	9
3.3.2	Concepts	10
3.3.3	Implications	11
3.3.4	Current State of Implementation	12
3.4	NFN	12
3.5	eXpressive Network Architecture	12
3.6	NovaGenesis	12
3.7	NEBULA	12
3.8	MobilityFirst	12
3.9	ChoiceNet	12
3.10	Comparison & Evaluation	12

4	Recursive InterNetwork Architecture	13
4.1	Premise	13
4.2	Concepts	13
4.2.1	Distributed Application	13
4.2.2	Error And Flow Control Protocol	13
4.2.3	Relaying and Multiplexing Task	13
4.3	Implications	13
4.4	Current State of Implementation	13
4.5	Evaluation	13
5	Implementation of Relaying and Multiplexing Task	14
5.1	OMNeT++	14
5.2	RINASim	14
5.3	Process of development	14
5.4	Implementation Design	15
5.4.1	Separation of mechanism and policy	15
5.4.2	Module structure	15
5.4.3	Module parameters	16
5.4.4	Module workflow	16
5.4.5	Module management	16
5.5	16
6	Testing and Evaluation	17
7	Conclusion & Future Development	18
A	CD Contents	19
B	Installation and Usage	20

Chapter 1

Introduction

(~2 ns)

1.1 Goals

The theoretical part of this thesis aims to describe alternatives to the currently prevalent network architectures. Since the Internet is by far the largest and most important example of an internetwork, its underlying architecture shall be used as a base for comparison. This is only fitting since nearly all of the network architecture research is directed towards improving the building blocks of the Internet.

The technical report describes design and implementation of a significant part of an OMNeT++ simulation model of one of the presented architectures, RINA.

1.2 Thesis Structure

Part Two describes the historical events that led to the current state of the Internet and hints on the shortcomings and weak parts of current technology which create the need for an alternative architecture research.

Part Three presents an outline of current research directions in field of network architecture and a brief overview of related research projects.

Part Four takes a closer look at Recursive InterNetwork Architecture, an IPC-based network architecture. It contains description of the architecture and its advantages.

Part Five describes implementation of RINA's Relaying and Multiplexing Task in OMNeT++.

Part Six presents evaluation of the implementation in form of test outputs.

Part Eight wraps it all up. ;=)

Chapter 2

Problems of the current Internet

(~5 ns)

The Internet could be, with no doubt, considered one of the most important technological achievements of the 20th century. It has brought a previously unimaginable degree of interconnection and information access to the whole world and its importance keeps growing even decades after its inception.

Nevertheless, the very basic core of its technology was constructed over three decades ago, back when the demands on internetworking capabilities were nowhere compared to the present situation. During the Internet's growth, whenever there has been a problem that required a solution (e.g. IPv4 address depletion), it's been usually dealt with in a non-intrusive evolutionary fashion by applying a new principle on top of the underlying technology. In another words, problems have been mostly solved by adding a new protocol to the TCP/IP protocol stack (e.g. RSVP for QoS guarantee).

This way of improving the Internet's base technology is only logical since each paradigm shift in foundations of the Internet can require a long and expensive transfer of existing network configurations to the new technology. The most notable example of this is the internet layer protocol IPv6 requiring explicit firmware support from active network components. The problem of IPv4 space exhaustion has been known of since the beginning of 1990s, first formal IPv6 specification arose in 1996 (RFC 2460) and the first IPv6 routers emerged in 2004 – and yet, as of 2015, two years after the top-level IPv4 pool exhaustion, IPv6 still represents only a miniscule fraction of total traffic on the Internet. For example, Google IPv6 adoption statistics indicate around 6% coverage amongst the users of its services.

As such, some of the Internet's problems are inherent because of the base design and it's usually difficult, if not impossible, to solve them incrementally. The following sections illustrate several of those problems.

2.1 Node naming

TODO: explain Saltzer's proposal and why it isn't used

Thus, both MAC addresses and IP addresses serve as point-of-attachment addresses and effectively locate the same thing, which is a host interface. This implies that the Internet effectively forwards PDUs on interface addresses, which has a great impact on difficulty of mobility and multihoming (both described in another sections). Borrowing a comparison from John Day's *Patterns in Network Architecture*, the lack of implicit logical addressing in the Internet “feels like accessing memory with DOS instead of Unix”.

2.2 Lack of Multihoming

Since IP addresses serve as point-of-attachment addresses (i.e. one per each computing system interface) and routing is done exclusively on IP's layer, there isn't any implicit mechanism for distinguishing whether multiple IP addresses belong to a single node. The insufficient base for multihoming support is also one of the oldest recognized problems of the Internet: it has become apparent back in 1972, when Tinker Air Force Base joined ARPANET and voiced a request for redundant connections to a single node to ensure reliability. In spite of this, a switch to a new protocol suite that happened 11 years later (on "flag day") didn't bring any solution to this problem.

Since then, some attempts were made to implement multihoming on top of the current architecture.

2.2.1 SCTP

Message-oriented transport protocol SCTP provides a partial support. Two SCTP hosts are able to provide each other with lists of alternate IP addresses that can be used as fallback points of attachment of the same application in case of failure of the primary address. During such session, each SCTP host needs to continually check other host's endpoints by heartbeat packets to make sure they are accessible via other sessions. However, thus far, multiple reasons have been preventing SCTP from becoming a widely known and used solution; the biggest disadvantage lies in the fact that due to TCP/IP not having a distinct session layer on top of the network layer (such as in ISO/OSI stack), the transport protocol must be explicitly specified by the application using sockets API. Therefore, SCTP adoption would require a rewrite of network-aware applications themselves. Other SCTP adoption issues include operating system support (Microsoft Windows requires a third-party kernel driver) and weak awareness of its existence outside networking community.

2.2.2 BGP Multihoming

Another implementation of multihoming capabilities can be seen in Border Gateway Protocol, which provides means for load-balancing and fallback over multiple links. To make use of such multihoming over the Internet, a public IP address range and an Autonomous System number are required. BGP Multihoming is one of the most significant causes contributing to the growth of Internet routing table.

2.2.3 Multipath TCP

The most recent TCP/IP multihoming initiative is a TCP extension Multipath TCP (MPTCP) which is currently in its experimental phase, although a large scale commercial deployment was already made by Apple for its Siri network application in mobile operating system iOS 7.

2.3 Lack of Mobility

Since a host location is determined by its IP address and IP addressing is geographically dependent, mobility with re-locatable addresses is essentially non-existent.

TODO: just a placeholder, re-read the relevant chapter of Patterns

2.4 Lack of Multicast

TODO: better get well informed on this one first

2.5 Weak security


The specifications of the fundamental protocols of TCP/IP stack – IP, TCP and DHCP – were originally finished at the beginning of 1980s. Computer internetworking in pre-Internet era was still a matter of closed research exercised in small scientific communities, so most of the design effort was given to making things work and network security wasn't the main aim. However, the Internet has since then turned into a massive world-wide internetwork connecting people of different types and agendas. Naturally, once the Internet began to be used for transferring sensitive data (especially by companies), cyber crime started to emerge as well and some attention was turned to security aspects of Internet protocol (or lack thereof).

Over the decades, a large amount of security flaws has been discovered and continually exploited. Some examples of these are Denial-of-service attacks (ICMP flood, SYN flood, CAM flood), Man-in-the-Middle attacks, IP address spoofing, DHCP spoofing, ARP spoofing and ARP cache poisoning. The requirement of every secure network is to carefully mitigate all of them.

2.6 Router Table Size Growth

Default-free zone (DFZ) is the collection of all Internet autonomous systems (AS) that do not require a default route to route a packet to any destination. Since they comprise the root of the Internet's routing infrastructure, their database must be complete.

With the increasing number of hosts connected to the Internet, the DFZ routing table sizes grow as well.



While the exponential growth observed during the 1990s was later mitigated by mass deployment of CIDR and route aggregation, the number of items is still increasing super-linearly and the high-end router hardware needs to keep up, especially with increasing use of BGP-based multihoming and IPv6. This can sometimes lead to scalability problems, as

in August of 2014 where reaching the 512k entry limit of multiple routers caused globally observable outages.

As of year 2015, the Internet routing table is consisted of over 560k entries, which requires over 0.5 GB of memory on each router. It's been commonly believed that Moore's law will ensure that the technology of high-end routers will keep scaling along with the increased demands, but recent research proves otherwise.

Chapter 3

Overview of Alternative Architectures

(~12 ns)

This section gives an overview of paths that were pursued in the field of network architecture research.

3.1 Design Approaches

Over the past several years, the networking research community has exhibited many attempts of moving the field forward. The undertaken research directions are often classified into one of two groups:

- evolutionary design: backward-compatible solutions that are incrementally deployable on top of the current Internet (e.g. SDN, LISP, DiffServ, IntServ), or
- clean slate design: designing completely new standalone architectures that aren't constrained by Internet technology's limitations (e.g. RINA, NDN)

Considering the scope of this thesis, the focus will be given exclusively to “clean-slate design” architectures.

3.2 Research Initiatives

3.2.1 NewArch Project

3.2.2 FIRE

3.2.3 FIBRE

3.2.4 Future Internet Architecture


In 2010, National Science Foundation funded five projects as a part of this program: Named Data Networking, MobilityFirst, NEBULA, eXpressive Network Architecture and ChoiceNet.

3.3 Named Data Networking

(~8 ns)

3.3.1 Premise


As described in section 2.1, the current Internet has its roots in telecommunication technology as this was the only area from which any inspiration could be taken. Because of this, technology of the Internet has been built on the paradigm of host-to-host communication. This can be visually presented on a “hourglass model”, which indicates that while there’s a wide array of technologies in use in the lower and higher layers, the hourglass’s thin waist of end-to-end communication is the key static part binding different networks together.



Building the Internet as a communication network was an obvious choice, especially since most of communication of the early ARPANET consisted of connecting terminals to mainframes and executing remote procedures on them.

However, while the basic paradigm of host-to-host communication hasn’t changed, the way we use the Internet has gone in an entirely different direction: the Internet has become mostly a content distribution network. Since the mechanism of communication over the Internet is based on creating and maintaining end-to-end connections, this creates an enormous amount of data redundancy.

Named Data Networking proposes a solution for the problem: instead of working with the source/destination node identifiers, the “thin waist” of the Internet should work with names of data chunks.

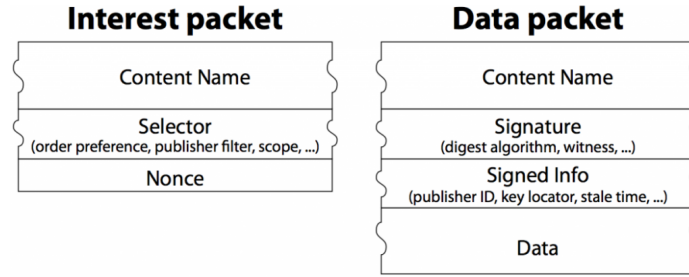


3.3.2 Concepts

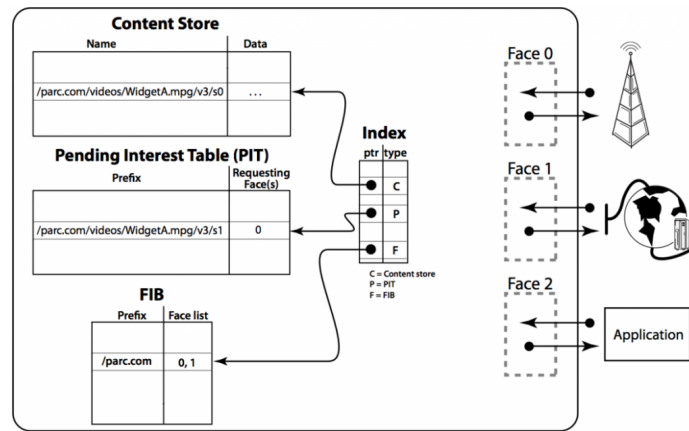
The Building Blocks

The NDN architecture specifies:

- two types of packets: an interest packet (containing the name of desired data) and a data packet (containing the requested data)



- two types of hosts: a consumer (data requester) and a producer (data provider)
- a router maintaining three fundamental data structures:
 - Forwarding Information Base (forwarding table)
 - Pending Interest Table (maintaining currently active requests)
 - Content Store (data cache)



The Communication Model

Communication in NDN is driven by the data receiver, i.e. consumer.

1. The consumer sends out an “interest packet” containing the name of the desired data.
2. When a router receives the interest packet, it first consults its Content Store for requested data.
 - If the data requested by the interest packet are present, they are returned in the direction of the requesting interface.
 - Otherwise, it'll look up the Pending Interest Table.

- If there's an entry present for the named data request, the entry is updated by adding the originating interface into the list of requesting interfaces, thus aggregating the new request together with an existing one.
 - Otherwise, a new entry is inserted, a FIB lookup is made and the interest packet is forwarded to interface(s) returned by the FIB.
3. A data packet is returned to the router by either the producer or another router with cached data. The router finds a matching PIT entry and forwards the data to all interfaces listed in the PIT entry. The PIT entry is then removed and data are cached into the Content Store.

3.3.3 Implications

Data Caching

The most significant feature of NDN is its native support for caching all sorts of data inside the network itself. While this should be beneficial mostly for static data such as web pages and images, dynamic content could take advantage of this as well in case of multicasting or packet retransmission on packet loss.

Routing and Forwarding

NDN routes and forwards packets on names, which eliminates four problems caused by address forwarding in the IP architecture: address space exhaustion, NAT traversal, mobility, and address management.

- There is no address exhaustion problem since the namespace is unbounded.
- There is no NAT traversal problem since a host does not need to expose its address in order to offer content.
- Mobility, which requires changing addresses in IP, no longer breaks communication since data names remain the same.
- Finally, address assignment and management is no longer required in local networks, which is especially empowering for embedded sensor networks.

Since the forwarding mechanism otherwise bears a strong resemblance to the forwarding mechanism in IP networks, lot of the existing research on IP routing could be applied to NDN as well: protocols like BGP, IS-IS or OSPF can be, for the most part, reused with minor modifications (e.g. routers would announce name prefixes instead of IP prefixes).

Security

While TCP/IP with its end-to-end communication paradigm relies on setting up a secure channel between two hosts and transmitting data through it, NDN requires each data chunk to be signed together with its name. Thus, since security is built into the data itself, there's no need for a direct host-to-host secure channel to be created.

3.3.4 Current State of Implementation

3.4 NFN

NFN's founders present it as "a companion project to Named Data Networking". While CCN (and, in turn, NDN) is based on the idea of resolving names to content, NFN decides to take an even more general approach: resolving names to computations.

3.5 eXpressive Network Architecture

As presented in the previous sections, some of the future architecture research is centered around the idea of replacing Internet's narrow waist of end-to-end communication with a different principal (e.g. NDN and its named content). XNA takes this approach one step further and builds on one basic premise: the narrow waist should provide support for multiple principals and the ability to evolve its functionality to accomodate new principals over time.

3.6 NovaGenesis

3.7 NEBULA

3.8 MobilityFirst

3.9 ChoiceNet

3.10 Comparison & Evaluation

some nice tables + a paragraph or two

Chapter 4

Recursive InterNetwork Architecture

(~8 ns) In 2008, computer scientist John D. Day has published a book that marked the culmination of his long-time goal of rediscovering the way we think about computer networks. The book was called Patterns in Network Architecture and in it, Day singlehandedly proposed a clean-slate approach to computer architecture that aims to get rid of most of TCP/IP's drawbacks.

The book can be roughly divided into two separate parts: in the first part, Day attempts to decompose mechanisms used in TCP/IP to their basic parts and put them into historical and socio-economical context. He eventually discovers that the currently prevalent layered approach to network architecture is needlessly complex, because each layer consists of the same mechanisms acting on a different scope.

The second part takes into account all of the facts discovered in the first part and uses them as foundations to build an entirely new concept of network architecture from scratch: the recursive IPC model (originally Network IPC Architecture, NIPCA).

RINA (Recursive InterNetwork Architecture) is a currently researched network architecture based on fundamental principles described by Day.

4.1 Premise

4.2 Concepts

4.2.1 Distributed Application

4.2.2 Error And Flow Control Protocol

4.2.3 Relaying and Multiplexing Task

4.3 Implications

4.4 Current State of Implementation

4.5 Evaluation

Chapter 5

Implementation of Relaying and Multiplexing Task

(~10 ns)

5.1 OMNeT++

OMNeT++ is an open-source discrete simulation framework used primarily in the field of network simulation. In this context, the term “network” refers to the more general meaning of the word, which means the framework can be used not only for simulation of TCP/IP networks (especially in conjunction with the INET library), but it also provides means for simulating other networked systems such as on-chip networks or queuing networks. As we’re implementing a clean-slate architecture from the ground up, this is an ideal approach.

OMNeT++ provides a component architecture for models. Components (modules) are programmed in C++, then assembled into larger components and models using the high-level language NED. In theory, there are no limits for networks modelled by NED and the only constraint is given by a computing platform processing power.

5.2 RINASim

RINASim, developed by networking research group at Faculty of Information Technology of Brno University of technology, is an OMNeT++ library developed for project PRISTINE. The purpose of the library is to provide a framework for building RINA networks. RINASim is used primarily by other PRISTINE researchers to experiment with the architecture and efficiently evaluate their working theories.

5.3 Process of development

Since RINA is still an emerging architecture, some of its parts (e.g. SDU Protection) are still being actively worked on and the current implementations aren’t yet fit for production purposes. Because of this, most of development was driven by PRISTINE researchers’ feature requests and frequent mutual feedback regarding architectural or implementational issues.

5.4 Implementation Design

In RINASim, all functionality of RMT including a policy architecture is encompassed in a single compound module named “RelayAndMux” which is present in every IPC process. The module serves for (de)multiplexing, relaying and aggregating PDUs of data flows traversing the IPC processes.

5.4.1 Separation of mechanism and policy

Some aspects of RMT functionality such as queue scheduling and performance monitoring are, by their nature, suited for separation of mechanism and policy. In another words, we need to take into account that architecture implementators will probably want to specify custom behavior for some of the logic present in RMT.

For this purpose, several algorithms used by RMT are meant to be user-configurable:

- RMT scheduling policy: The scheduling algorithm that determines the order input and output buffers are serviced (e.g. by number of waiting PDUs in buffers or with fair queuing).
- RMT monitoring policy: A queue monitoring algorithm that is invoked each time a PDU enters or leaves a queue. This policy should compute variables (such as average queue length) to be used in decision process of other policies.
- RMT maxqueue policy: An algorithm that is invoked each time a number of PDUs waiting in a queue exceeds the queue’s threshold. This policy is mostly used for congestion control mechanisms (e.g. by dropping or marking the last PDU in a queue).

5.4.2 Module structure

(the nasty guts of RelayAndMux.png)

RelayAndMux consists of multiple simple modules of various types, some of which are instantiated only dynamically at runtime.

Static modules:

- RMT, the central logic of Relaying And Multiplexing task that decides what should be done with messages passing through the module.
- RMTModuleAllocator, a control unit providing an API for adding and deleting instances of dynamic modules (RMTQueue, RMTPort).
- SchedulingPolicy, a scheduling policy of an RMT instance.
- MonitorPolicy, a monitoring policy of an RMT instance.
- MaxQPolicy, a maxqueue policy of an RMT instance.

Dynamic modules:

- RMTPort, a representation of one endpoint of an (N-1)-flow.
- RMTQueue, a representation of either input or output queue. The number of RMTQueues per RMTPort is determined by Resource Allocator policies.
- RMTPortWrapper, a compound module encapsulating an (N-1)-port and its queues.

5.4.3 Module parameters

Some of the modules contain user-configurable parameters that can be used to alter the modules' behavior. The whole list of parameters can be found in attachment A2.

5.4.4 Module workflow

diagrams, descriptions, pie charts etc.

5.4.5 Module management

RMT's purpose in an IPC process is fairly straightforward: providing a stateless function for relaying PDUs to their predetermined destinations and multiplexing PDUs of multiple data flows onto a common predetermined medium. The entire management of the RMT is decoupled and exercised by some of the functions of the Resource Allocator. This includes tasks such as:

- Initial setup of RMT mode of function. When a RMT instance is located inside a bottommost IPC process that doesn't work with any further (N-1)-IPC processes, the RMT is switched to an "onWire" mode that functions over a simple medium instead of an IPC connection.
- Forwarding table management. The content of the PDU forwarding table is generated by the PDU Forwarding Table Generator module which generally exchanges routing information with other IPC processes by some sort of routing protocol.
- Queue allocation. PDUs traversing an (N-1)-port may be momentarily buffered in input or output queues; the number of input and output queues per (N-1)-port and assignment of traffic classes to queues (e.g. all-in-one or fair queuing) is determined by Resource Allocator policies.
- (N-1)-port creation & removal. Since Resource Allocator manages (N-1)-flows with other IPC processes, it also provides (N-1)-ports (or handles) for RMT.
- (N-1)-port data rate control. In some scenarios, it may be required for an (N-1)-port to cease/slow down sending or providing more data because of congestion. Resource Allocator can momentarily disable or slow down data rate on distinct ports if this is required by EFCP instances.

5.5

Chapter 6

Testing and Evaluation

(~4 ns)

Chapter 7

Conclusion & Future Development

(~1 ns)

Appendix A

CD Contents

Appendix B

Installation and Usage